

Taiyu Zhang

xiannian0311@gmail.com · +8618843980900 · [ssstirm.github.io](https://github.com/ssstirm)



SUMMARY

PhD candidate at KU Leuven Computer Science, researching 3D multimodal intelligent agents in immersive environments. B.S. from USTC's School of Information; M.S. from the Department of Electronic Engineering, Tsinghua University. Research focuses primarily on Human-Computer Interaction and generative foundation models. Project experience covers AI Agents, XR interaction, Language model generation, Full-body motion generation, and World models.

EXPERIENCE

Part-time Research Engineer — Moxin Technology

Mar 2026 – May 2026

- On Huawei Ascend servers, performed self-forcing distillation from a 14B teacher to a 14B student; added camera-trajectory control conditioning on top.
- Applied 4-step few-step distillation to the student on a 1000-frame long-video pretrained model.

NLP Researcher — iFlyTek

2021 – 2022

- Multiple national-level R&D projects; main owner of the multilingual semantic-understanding component of the iFlyTek Digital Human — data curation, model training, and evaluation.
- Deployed in the 2022 Beijing Winter Olympics and other state-level events.
- Long-sentence-completion GPT pretraining for the iFlyTek input method (see Selected Projects).

SELECTED PROJECTS

iFlyTek IME Long-Sentence Completion — *Lead researcher, GPT pretraining from scratch*

2021 – 2022

- Built a 400 GB → 40 GB pretraining-data pipeline: deduplication, language and quality filtering, topic balancing, formatting normalization.
- Pretrained a GPT-style language model from scratch on the cleaned 40 GB Chinese corpus.
- Shipped in the iFlyTek input method, powering long-text completion and poetry-creation features in production.

LingBot Distillation — *Lead researcher, 14B video diffusion → 4-step student on Ascend*

2026.03 – 2026.05

- Reproduced LingBot's distillation on Huawei Ascend (CANN); teacher: camera-trajectory control + 1000-frame long-video coverage. 40–50-step teacher → 4-step student; > 50-min generation → < 4 min.
- Resolved self-forcing memory bottleneck with two-pass forward + activation recompute (gradient-accumulation style); proposed a new ODE-initialization that uses training data to init student directly from teacher, removing explicit ODE-pair generation and stabilising distillation.

Claw-AI-Lab — *Co-developer*

2025 – Present

- Multi-agent autonomous research platform: one prompt → paper, code, figures, and experiment logs. 1.3k+ GitHub stars.
- Contributed across the multi-agent scheduler, the local-execution harness, and the real-time monitoring dashboard.

3D Character Animation Generation — *Lead researcher, autoregressive token models + distillation*

2023 – 2025

- Tracked both branches of text-to-motion: discrete-token autoregressive (T2M-GPT, Momask, MotionGPT) and diffusion-based (MDM and successors).
- Swapped VQ-VAE for VQGAN as the motion tokenizer (improved sequence reconstruction); applied few-step distillation to a diffusion motion generator (preserved FID).

Focus Agent — *First author, IVA 2024*

2023 – 2024

- LLM-driven virtual focus group: each participant is an agent with persona, reasoning style, and memory; moderator steers toward an interview goal.
- Published at the 24th ACM International Conference on Intelligent Virtual Agents.

AI Game NPC — *First author, VR + LLM agents*

2024 – 2025

- LLM-driven intelligent NPC system with a three-layer architecture: Whisper-based speech perception, GPT-4 cognition with direct game-engine state queries, and TTS-plus-action executor; deployed on Apple Vision Pro for evaluation.
- 24-participant within-subjects study across four assistance levels; found that assistance reduces workload but full autonomy hurts UX, with four distinct player types showing different interaction preferences.

Coembodiment AI — *First author*

2025 – 2026

- First validation of users' subjective experience under human-AI coembodiment of one virtual body; two-study programme — Study 1 surfaces the problem, Study 2 proposes a method.
- A non-linear adaptive control function delivers substantially better UX than gate / linear baselines; result generalises to future human-robot collaborative control.

EDUCATION

Ph.D., Computer Science — KU Leuven

2022 – Present

Research direction: 3D multimodal intelligent agents in immersive environments. Advisor: Prof. Adalberto L. Simeone.

M.S., Electronic Engineering — Tsinghua University

2018 – 2021

Thesis: ML-based quantitative analysis of soil elements via LIBS. Advisor: Prof. Xiaohong Ma.

B.S., Electronic Engineering & Information Science — USTC

2014 – 2018

Thesis: SRCNN-based image super-resolution (Outstanding Undergraduate Thesis). Advisor: Prof. Jian Tang.

SELECTED PUBLICATIONS

Focus Agent: LLM-Powered Virtual Focus Group.

Zhang, T., et al. *Proceedings of the 24th ACM International Conference on Intelligent Virtual Agents (IVA)*, 2024.